

Dynamic Programming

Introduction

- ▶ The imperialism of Dynamic Programming (Ljungqvist & Sargent)

Markov chain and Markov process

- ▶ Stochastic process: family of random variables indexed by time
- ▶ A stochastic process has the Markov property if its future evolution depends only on its current state.
- ▶ Special cases:

	Discrete States	Continuous States
Discrete Time	Discrete Markov Chain	Continuous Markov Chain
Continuous Time	Markov Jump Process	Markov Process

Stochastic matrices

- ▶ a $n \times n$ matrix is said to be **stochastic** if all the lines sum to 1
- ▶ let's a vector $\mu_t \in R^n$ with positive components denote a distribution of mass between n different states
- ▶ if $\sum_{i=1}^n \mu_{i,t} = 1$ then μ_t is a probability density

Simulation

- ▶ Consider: $\mu'_{i,t+1} = \mu'_t P$
- ▶ We have $\mu_{i,t+1} = \sum_{k=1}^n \mu_{k,t} P_{k,i}$
- ▶ And: $\sum_i \mu_{i,t+1} = \sum_i \mu_{i,t}$
- ▶ Postmultiplication by a stochastic matrix preserves the mass.
- ▶ Interpretation: P_{ij} is the fraction of the mass initially in state i which ends up in j

Example

$$\underbrace{\begin{pmatrix} ? & ? & ? \end{pmatrix}}_{\mu'_{t+1}} = \underbrace{\begin{pmatrix} 0.5 & 0.3 & 0.2 \end{pmatrix}}_{\mu'_t} \begin{pmatrix} 0.4 & 0.6 & 0.0 \\ 0.2 & 0.5 & 0.3 \\ 0 & 0 & 1.0 \end{pmatrix}$$

Representation as a graph

```
{% dot attack_plan.svg digraph G { rankdir=LR A B C A -> B  
[label=0.4] A -> C [label=0.6] C -> C [label=1.0] B -> A  
[label=0.2] B -> B [label=0.5] B -> C [label=0.5] } %}
```

Probabilistic interpretation

- ▶ Denote by $S = (s_1, \dots, s_n)$ a finite set with n elements ($|S| = n$).
- ▶ A **Markov Chain** with values in S and with transitions given by a stochastic matrix $P \in R^n \times R^n$ is a *stochastic process* $(X_t)_{t \geq 0}$ such that

$$P_{ij} = \text{Prob}(X_{t+1} = s_j | X_t = s_i)$$

- ▶ In words, line i describes the conditional distribution of X_{t+1} conditional on $X_t = s_i$.

What about longer horizons?

- ▶ It is easy to show that for any k , P^k is a stochastic matrix.
- ▶ P_{ij}^k denotes the probability of ending in j , after k periods, starting from i
- ▶ Given an initial distribution $\mu_0 \in R^n$
 - ▶ Which states will be visited with positive probability between $t=0$ and $t=k$?
 - ▶ What happens in the very long run?
- ▶ We need to study a little bit the properties of Markov Chains

Connectivity

- ▶ Two states s_i and s_j are connected if $P_{ij} > 0$
- ▶ We call incidence matrix: $\mathcal{I}(P) = (\delta_{P_{ij}>0})_{ij}$
- ▶ Two states i and j communicate with each other if there are k and l such that: $(P^k)_{i,j} > 0$ and $(P^l)_{j,i} > 0$
 - ▶ it is an equivalence relation
 - ▶ we can define equivalence classes
- ▶ A stochastic matrix P is irreducible if all states communicate
 - ▶ there is a unique communication class

Connectivity and irreducibility (example from QE)

Connectivity and irreducibility (example from QE)

Aperiodicity

- ▶ Are there cycles? Starting from a state i , how long does it take to return to i ?
- ▶ The **period** of a state is defined as

$$\gcd(k \geq 1 | (P^k)_{i,i} > 0)$$

- ▶ If a state has a period $d > 1$ the chain returns to the state only at dates multiple of d .

Aperiodicity

example

Stationary distribution

- ▶ μ is a **stationary** distribution if $\mu' = \mu'P$
- ▶ Theorem: there always exists such a distribution
 - ▶ proof: Brouwer theorem
 - ▶ $f : \mu \rightarrow (\mu'P)'$
- ▶ Theorem:
 1. if P is irreducible the fixed point μ^* is unique
 2. if P is irreducible and aperiodic $|\mu_0'P^k - \mu^*| \xrightarrow[k \rightarrow +\infty]{} 0$ for any initial distribution μ_0
- ▶ We then say the Markov chain is **ergodic**.

Stationary distribution (proof)

- **Brower's theorem:** Let \mathcal{C} be a compact convex subset of R^n and f a continuous mapping $\mathcal{C} \rightarrow \mathcal{C}$. Then there exists a fixed point $x_0 \in \mathcal{C}$ such that $f(x_0) = x_0$

Stationary distribution?

How do we compute the stationary distribution?

- ▶ Simulation
- ▶ Linear algebra
- ▶ Decomposition

Simulating a Markov Chain

- ▶ Very simple idea: start with μ_0 and compute the iterates recursively
 - ▶ $\mu'_{n+1} = \mu'_n P$
 - ▶ convergence is linear

Using Linear Algebra

- ▶ Find the solution of $\mu'(P - I) = 0$?
 - ▶ not well defined, 0 is a solution
 - ▶ we need to incorporate the constraint $\text{sum}(\mu_i) = 1$
- ▶ Define $M_{ij} = (P - I)_{ij}$ if $j > 1$, 1 if $j = 1$
- ▶ Look for a solution μ of $\mu'M = (\delta_{i=1})$ with a linear algebra solver
 - ▶ if the solution completes, there is a unique solution

Further comments

- ▶ Knowledge about the structure of the Markov Chain can help speedup the calculations
- ▶ There are methods for potentially very-large linear system
 - ▶ Newton-Krylov based methods, GMRES
- ▶ Basic algorithms are easy to implement by hand
- ▶ QuantEcon toolbox has very good methods to study markov chains

Dynamic Programming: notations

General Formulation

Markov Decision Problem

- ▶ states: $s \in S$
- ▶ actions: $x \in X(s)$
- ▶ transitions: $\pi(s'|s, x)$
 - ▶ probability of going to s' in state s , given action x

Objective (finite horizon)

- ▶ policy: $x() : s \rightarrow x \in S(x)$
 - ▶ deterministic policy
 - ▶ given $x()$, the evolution of s' is a Markov process.
- ▶ reward: $r(s, x)$
 - ▶ felicity, intratemporal utility
- ▶ expected lifetime reward: starting from s
 - ▶ $R(s; x()) = E_0 \sum_t^T \delta^t [r_t]$
 - ▶ $\delta \in [0, 1[$: discount factor, $T \in \{N, \infty\}$: horizon

Classes of Dynamic Optimization

The formulation so far is very general. It encompasses several variants of the problem:

- ▶ finite horizon vs infinite horizon
- ▶ discrete-space problem vs continuous-state space problem

There are also variants not included: - non time-separable problems - non time homogenous problem - learning problems (bayesian updating, reinforcement learning)

Finite horizon vs infinite horizon

- ▶ Recall objective: $V(s; x()) = \max E_0 \sum_{t=0}^T \delta^t [r(s_t, x_t)]$
- ▶ If $T < \infty$, the decision in the last periods, will be different from before
 - ▶ one must find a decision rule $\pi_t()$ per period
 - ▶ or add t to the state space: $\tilde{S} = S \times [0, T]$
- ▶ If $T = \infty$, the continuation value of being in state s_t is independent from t

$$\begin{aligned} V(s; x()) &= E_0 \max \sum_{t=0}^{T_0} \delta^t [r(s_t, x_t)] + \delta^{T_0} E_0 \sum_{t=T_0}^{\infty} \delta^t [r(s_t, x_t)] \\ &= E_0 \left[\max \sum_{t=0}^{T_0} \delta^t [r(s_t, x_t)] + \delta^{T_0} V(s_{T_0}; x()) \right] \end{aligned}$$

Continuous vs discrete

- ▶ State space S :
 - ▶ continuous: $\subset R^n$
 - ▶ discrete: $S = (s_1, \dots, s_n)$ (today)
- ▶ General approach is the same but implementation is different:
- ▶ Continuous problem:
 - ▶ $x(s)$, $V(s; \pi)$ require an infinite number of coefficients
 - ▶ one must discretize the initial problem and solve an approximate version
- ▶ Discrete problem:
 - ▶ there is a finite number of policies, they can be represented exactly
 - ▶ unless $|S|$ is very large (cf go game)

Non time separable

- ▶ For instance Epstein-Zin preferences:

$$\max V(; c())$$

where $V_t = (1 - \delta) \frac{c_t^{1-\sigma}}{1-\sigma} + \delta [E_t V_{t+1}^\alpha]^\frac{1}{\alpha}$

- ▶ Why would you do that?
 - ▶ to disentangle risk aversion and elasticity of intertemporal substitution
 - ▶ robust control
- ▶ You can still use ideas from Dynamic Programming.

Non homogenous preference

- Look at the $\alpha - \beta$ model.

$$V_t = \max \sum_t^{\infty} \beta_t U(c_t)$$

where $\delta_0 = 1$, $\delta_1 = \alpha$, $\delta_k = \alpha\beta^{k-1}$

- Makes the problem time-inconsistent: the optimal policy you would choose for the continuation value after T is not the same if you maximize it in expectation from 0 or at T .

Learning problems

- ▶ Bayesian learning: Uncertainty about some model parameters
 - ▶ ex: variance and return of a stock market
 - ▶ agent models this uncertainty as a distribution
 - ▶ agent updates his priors after observing the result of his actions
 - ▶ actions are taken optimally taken into account the revelation power of some actions
- ▶ Is it good?
 - ▶ clean: the *rational* thing to do with uncertainty
 - ▶ super hard: the state-space should contain all possible priors
 - ▶ mathematical cleanness comes with many assumptions
- ▶ Used to estimate rather big (mostly linear) models

Learning problems (2)

- ▶ Reinforcement learning
 - ▶ model can be partially or totally unknown
 - ▶ decision rule is updated by observing the reward from actions
 - ▶ no priors
 - ▶ solution does not derive directly from model
 - ▶ can be used to solve dynamic programming problems
- ▶ Good solutions maximize a criterion similar to lifetime reward but are usually not optimal:
 - ▶ usually evaluated by replaying the game many times
 - ▶ tradeoff exploration / exploitations

Examples

Examples (2)

Finite horizon DMDP

Finite horizon DMDP

When $T < \infty$. With discrete action the problem can be represented by a tree.

[Graph]

Finite horizon DMDP

- ▶ Intuition: backward induction.
 - ▶ Find optimal policy $x_T(s_T)$ in all terminal states s_T . Set $V_T(s_T)$ equal to $r(s_T, \pi_T)$
 - ▶ For each state $s_{k-1} \in S$ find $x_{k-1} \in X(s_{k-1})$ which maximizes

$$V_{k-1}(s_{k-1}) = \max_{x_{k-1}(s_{k-1}) \in X(s_{k-1})} r(s_{k-1}, x_{k-1}) + \underbrace{\delta \sum_{s_k \in S} p(s_k | s_{k-1}, x_{k-1}) V_k(s_k)}_{\text{expected continuation value}}$$

- ▶ Policies $x_0(), \dots, x_T()$ are “Markov-perfect”: they maximize utility on all subsets of the “game”
 - ▶ also from $t=0$

Remarks

- ▶ Can we do better than this naive algorithm?
 - ▶ not really
 - ▶ but we can try to limit S to make the maximization step faster
 - ▶ exclude a priori some branches in the tree using knowledge of the problem

Infinite horizon DMDP

Infinite horizon DMDP

- ▶ Horizon is infinite:

$$V(s; x()) = \max E_0 \sum_{t=0}^{\infty} \delta^t r(s_t, x_t)$$

- ▶ Intuition:
 - ▶ let's consider the finite horizon version $T < \infty$ and $T \gg 1$
 - ▶ compute the solution, increase T until the solution doesn't change
 - ▶ in practice: take an initial guess for V_T then compute optimal V_{T-1} , V_{T-2} and so on, until convergence of the V s

Infinite horizon DMDP (2)

- ▶ This is possible, it's called *Successive Approximation* or *Value Function Iteration*
 - ▶ how fast does it converge? *linearly*
 - ▶ can we do better? yes, *quadratically*

Successive Approximation

- Consider the decomposition:

$$V(s; x()) = E_0 \sum_{t=0}^{\infty} \delta^t r(s_t, x_t) = E_0 \left[r(s, x(s)) + \sum_{t=1}^{\infty} \delta^t r(s_t, x_t) \right]$$

or

$$V(s; x()) = r(s, x(s)) + \delta \sum_{s'} p(s'|s, x(s)) V(s'; x())$$

Successive Approximation (2)

- ▶ Taking continuation value as given we can certainly improve the value in every state \tilde{V} by choosing $\tilde{x}()$ so as to maximize

$$\tilde{V}(s; x(), \tilde{x}()) = r(s, \tilde{x}(s)) + \delta \sum_{s'} \pi(s'|s, \tilde{x}(s)) V(s'; x())$$

- ▶ By construction: $\forall s, \tilde{V}(s, \tilde{x}(), x()) > V(s, x())$
 - ▶ it is an “improvement step”
- ▶ Can $V(s, \tilde{x}())$ be worse for some states than $V(s, x())$?
 - ▶ actually no

Bellman equation

Idea: it should not be possible to improve upon the optimal solution. Hence the optimal value V and policy x^* should satisfy:

$$\forall s \in S, V(s) = \max_{y(s)} r(s, y(s)) + \delta \sum_{s' \in S} \pi(s'|s, y(s)) V(s')$$

with the maximum attained at $x(s)$. This is referred to as the **Bellman equation**.

Conversely, it is possible to show that a solution to the Bellman equation is also an optimal solution to the initial problem.

Bellman operator

- ▶ The function
$$G : V \rightarrow \max_{y(s)} r(s, y(s)) + \delta \sum_{s' \in S} \pi(s'|s, y(s)) V(s')$$
is known as the **Bellman operator**.
- ▶ Optimal value if a fixed point of G
- ▶ Can we show it's a contraction mapping ?

Blackwell's theorem

- ▶ Let $X \subset R^n$ and let $\mathcal{C}(X)$ be a space of bounded functions $f : X \rightarrow R$, with the sup-metric. $B : \mathcal{C}(X) \rightarrow \mathcal{C}(X)$ be an operator satisfying two conditions:
 1. (monotonicity) if $f, g \in \mathcal{C}(X)$ and $\forall x \in X, f(x) \leq g(x)$ then $\forall x \in X (Bf)(x) \leq (Bg)(x)$
 2. (discounting) there exists some $\delta \in]0, 1[$ such that:
$$B.(f + a)(x) \leq (B.f)(x) + \delta a, \forall f \in \mathcal{C}(X), a \geq 0, x \in X$$
- ▶ Then B is a contraction mapping with modulus δ .

Successive Approximation

- ▶ Using the Blackwell's theorem, we can prove the Bellman operator is a contraction mapping (do it).
- ▶ This justifies the Value Function Iteration algorithm:
 - ▶ choose an initial V_0
 - ▶ given V_n compute $V_{n+1} = G(V_n)$
 - ▶ iterate until $|V_{n+1} - V_n| \leq \eta$
- ▶ Policy rule is deduced from V as the maximand in the Bellman step

Successive Approximation (2)

- ▶ Not that convergence of V_n is geometric
- ▶ But x_n converges in finite time (X is finite)
 - ▶ surely the latest iterations are suboptimal
 - ▶ they serve only to evaluate the value of x^*
- ▶ In fact:
 - ▶ V_n is never the value of $x_n()$
 - ▶ should we try to keep both in sync?

Policy iteration for DMDP

- ▶ Choose initial policy $x_0()$
- ▶ Given initial guess $x_n()$
 - ▶ compute the value function $V_n = V(; x_n)$ which satisfies $\forall s, V_n(s) = r(s, x_n(s)) + \delta \sum_{s'} \pi(s'|s, x_n(s)) V_n(s')$
 - ▶ improve policy by maximizing in $x_n()$

$$\max_{x_n()} r(s, x_n(s)) + \delta \sum_{s' \in S} \pi(s'|s, x_n(s)) V_{n-1}(s')$$

- ▶ Repeat until convergence, i.e. $x_n = x_{n+1}$
- ▶ One can show the speed of convergence (for V_n) is *quadratic*
 - ▶ it corresponds to the Newton-Raphson steps applied to $V \rightarrow G(V) - V$

How do we compute the value of a policy?

- ▶ Given x_n , goal is to find $V_n(s)$ in
$$\forall s, V_n(s) = r(s, x_n(s)) + \delta \sum_{s'} \pi(s'|s, x_n(s)) V_n(s')$$
- ▶ Two(three) approaches:
 1. simulate the policy rule and compute $E [\sum_t \delta^t r(s_t, x_t)]$ with Monte-Carlo draws
 2. successive approximation: - put V_k in the rhs and recompute the lhs V_{k+1} , replace V_k by V_{k+1} and iterate until convergence
 3. solve a linear system in V_n
- ▶ For 2 and 3 it helps representing a linear operator M such that
$$V_{n+1} = R_n + \delta M_n \cdot V_n$$

Another approach consist: compute $(\sum_{t \geq 0} \delta^t M_n^t)$